

Workshop: Introduction to Interpretability in ML

Andras Zsom

Assistant Professor of the Practice of Data Science
Data Science Initiative, Brown University

Materials available at <https://github.com/azsom/ODSC-East-2022>

About me

- Born and raised in Hungary
- Astrophysics PhD at MPA, Heidelberg, Germany
- Postdoctoral researcher at MIT (still in astrophysics at the time)
- Started at Brown in 2015 as a Data Scientist in CCV (Center for Computation and Visualization)
- Lead Data Scientist since 2017 at CCV
- Assistant Professor since summer of 2021 at the DSI

Data Science at AI+Training

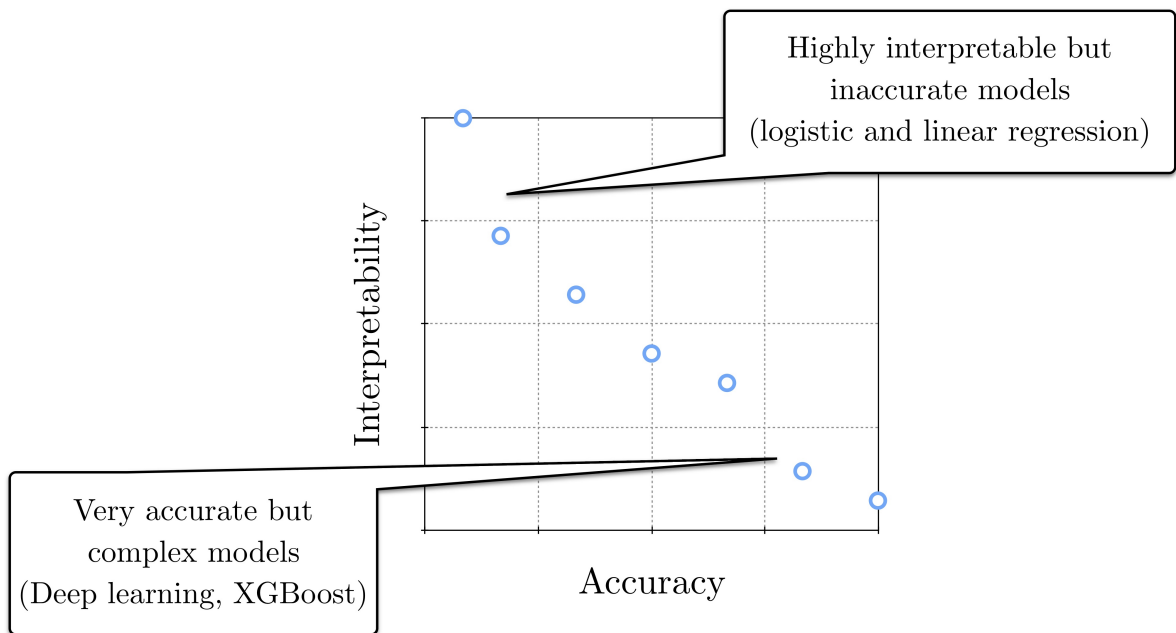
- Supervised Machine Learning Course Series - <https://app.aiplus.training/courses/supervised-machine-learning-series>
- 6 courses that walk through the main steps of developing an ML pipeline
- github repo available here: <https://github.com/azsom/Supervised-Learning>
- Week 7 of the ODSC ML Certification - <https://aiplus.training/certificates/>
- this workshop is based on course 6

Introduction to Interpretability in ML

By the end of this workshop, you will be able to

- calculate global explanations using the coefficients of linear models
- perform permutation importance using any ML model
- use the SHAP package to calculate local explanations

Motivation: simple vs. complex models



- local feature importance improves the interpretability of complex models
- check out [this page](#) for a good example

Motivation: build trust in your model

- debugging ML models is tough
 - a model that runs without errors/warning is not necessarily correct
- model seems accurate but does it make predictions the way you expect it to be?
 - wolf vs. husky classifier
- model inspection is important
 - feature importance metrics measure useful they are at predicting the target variable
 - global: one score per feature, a vector of shape $(1, n_features)$
 - general overview over a larger sample size
 - local: one score per feature and point, an array of shape $(n_samples, n_features)$
 - explanation specific to each sample
- discuss the interpretability results with a subject matter expert
 - if they say something is surprising, you either
 - have a bug in your code (more likely),
 - you discovered something new which is exciting!

Introduction to Interpretability in ML

By the end of this workshop, you will be able to

- **calculate global explanations using the coefficients of linear models**
- perform permutation importance using any ML model

- use the SHAP package to calculate local explanations

Math recap

Linear regression:

$$f(X_i) = y'_i = \theta_0 + X_{i1}\theta_1 + X_{i2}\theta_2 + \dots = \theta_0 + \sum_{j=1}^m \theta_j X_{ij},$$

Logistic regression:

$$f(X_i) = y'_i = \frac{1}{1+e^{-z}},$$

$$z = \theta_0 + \sum_{j=1}^m \theta_j x_{ij}$$

Here, y'_i is the prediction of the model for sample i using the X_i feature vector, and θ are coefficients.

What to do if you want to use the coefficients of your linear model as explanations

- **Look at the absolute value of the coefficients!**
 - a feature is important, if its coefficient's magnitude is large
 - doesn't matter if the coefficient is positive or negative
- **Standardize ALL features!**
 - most practitioners standardize the continuous features
 - but practitioners often forget to standardize the one-hot encoded or ordinal encoded features which is a mistake!
 - all features need to have the same mean and standard deviation if you want to use the coefficients to measure importance!

Simple ML pipeline

- we will use the diabetes dataset in this workshop
- <https://www4.stat.ncsu.edu/~boos/var.select/diabetes.html>
- "Ten baseline variables, age, sex, body mass index, average blood pressure, and six blood serum measurements were obtained for each of $n = 442$ diabetes patients, as well as the response of interest, a quantitative measure of disease progression one year after baseline."

```
In [1]: # read in the dataset
import pandas as pd
```

```
df = pd.read_csv('data/diabetes.csv')
df.head()

y = df['Y']
X = df.loc[:, df.columns != 'Y']
print(y.head())
print(X.head())
```

```
0    151
1     75
2    141
3    206
4    135
Name: Y, dtype: int64
```

	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6
0	59	male	32.1	101.0	157	93.2	38.0	4.0	4.8598	87
1	48	female	21.6	87.0	183	103.2	70.0	3.0	3.8918	69
2	72	male	30.5	93.0	156	93.6	41.0	4.0	4.6728	85
3	24	female	25.3	84.0	198	131.4	40.0	5.0	4.8903	89
4	50	female	23.0	101.0	192	125.4	52.0	4.0	4.2905	80

```
In [2]: # our ML pipeline
def MLpipe_KFold_RMSE(X, y, preprocessor, ML_algo, param_grid, random_state):
    '''
    This function splits the data to other/test (80/20) and then applies KFold
    The RMSE is minimized in cross-validation.
    '''

    # split data to other/test 80/20, and the use KFold with 4 folds
    X_other, X_test, y_other, y_test = train_test_split(X, y, test_size=0.2, random_state=random_state,
                                                         cv=kf)
    kf = KFold(n_splits=4, shuffle=True, random_state=random_state)

    # the sklearn pipeline
    if preprocessor == None:
        # no preprocessing
        pipe = Pipeline(steps=[('regressor', ML_algo)])
    else:
        pipe = Pipeline(steps=[('preprocessor', preprocessor),
                               ('regressor', ML_algo)])

    # loop through the hyperparameter combinations or use GridSearchCV
    grid = GridSearchCV(pipe, param_grid=param_grid, scoring='r2',
                        cv=kf, return_train_score=True)
    # for each combination, calculate the train and validation scores using the
    grid.fit(X_other, y_other)
    # find which hyperparameter combination gives the best validation score
    print(grid.best_params_)

    # calculate the test score
    test_score = grid.score(X_test, y_test)
    print(test_score, r2_score(y_test, grid.predict(X_test)))

    return grid.best_estimator_, test_score, (X_other, y_other), (X_test, y_test)
```

```
In [3]: # let's try three different preprocessing techniques and collect the model coefficients
from sklearn.linear_model import Ridge
from sklearn.pipeline import Pipeline
from sklearn.model_selection import ParameterGrid
from sklearn.model_selection import train_test_split
```

```

from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer
from sklearn.metrics import r2_score
from sklearn.preprocessing import StandardScaler, OrdinalEncoder
from sklearn.compose import ColumnTransformer

param_grid = {'regressor__alpha': [1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2]}

num_fts = ['AGE', 'BMI', 'BP', 'S1', 'S2', 'S3', 'S4', 'S5', 'S6']
ord_fts = ['SEX']

# ordinal encoder
ordinal_transformer = Pipeline(steps=[
    ('ordinal', OrdinalEncoder())])
# standard scaler
numeric_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())])

#
# ordinal encode SEX only!
#
preprocessor1 = ColumnTransformer(remainder = 'passthrough',
    transformers=[
        ('ord', ordinal_transformer, ord_fts)])

model, test_score, train_set, test_set = MLpipe_KFold_RMSE(X, y, preprocessor1,
print(test_score)
coefs_no_prep = model['regressor'].coef_
print(coefs_no_prep)

#
# ordinal encode SEX, and standardize continuous features
#
preprocessor2 = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, num_fts),
        ('ord', ordinal_transformer, ord_fts)])

model, test_score, train_set, test_set = MLpipe_KFold_RMSE(X, y, preprocessor2,
print(test_score)
coefs_some_prep = model['regressor'].coef_
print(coefs_some_prep)

#
# ordinal encode SEX, standardize continuous features, and standarize all featu
#
preprocessor3 = Pipeline(steps=[
    ('prep2', preprocessor2),
    ('scaler', StandardScaler())])

model, test_score, train_set, test_set = MLpipe_KFold_RMSE(X, y, preprocessor3,
print(test_score)
coefs_prep = model['regressor'].coef_
print(coefs_prep)

```

```

{'regressor__alpha': 1.0}
0.4520945451686046 0.4520945451686046
0.4520945451686046
[-22.80861461  0.14269511  5.90541998  1.19867986 -1.07900835
  0.62662466  0.3774738  9.77013169 60.79394666  0.21396887]
{'regressor__alpha': 1.0}
0.45418697357725524 0.45418697357725524
0.45418697357725524
[ 1.79050083 25.75648091 16.71212465 -34.65085803 17.05080545
 3.3870545 11.71899321 31.39253301 2.45268588 -22.69228495]
{'regressor__alpha': 1.0}
0.4541465207069819 0.4541465207069819
0.4541465207069819
[ 1.80734179 25.73269892 16.73429974 -34.67195409 17.05307485
 3.36991411 11.76426044 31.3783838 2.45813922 -11.44818951]

```

```

In [4]: import matplotlib.pyplot as plt
import numpy as np
plt.rcParams.update({'font.size': 14})

def plot_coefs(coefs,title,color,feature_names):
    feature_names = np.array(feature_names)
    sorted_indcs = np.argsort(np.abs(coefs))
    plt.barh(np.arange(10),coefs[sorted_indcs])
    plt.yticks(np.arange(10),feature_names[sorted_indcs])
    plt.xlabel('coefficient')
    plt.title(title,color=color)
    return

plt.figure(figsize=(10,5))

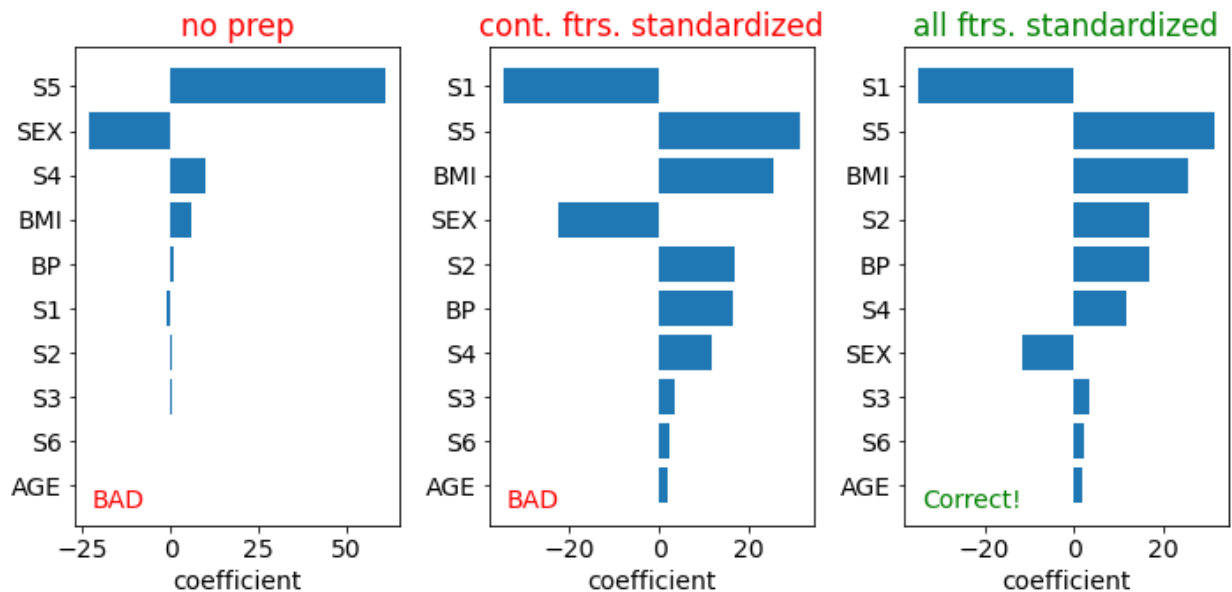
plt.subplot(1,3,1)
plot_coefs(coefs_no_prep,'no prep','r', ord_ftrs + num_ftrs)
plt.annotate('BAD',xy=(10,10),xycoords = 'axes points',color='r')

plt.subplot(1,3,2)
plot_coefs(coefs_some_prep,'cont. ftrs. standardized','r', num_ftrs + ord_ftrs)
plt.annotate('BAD',xy=(10,10),xycoords = 'axes points',color='r')

plt.subplot(1,3,3)
plot_coefs(coefs_prep,'all ftrs. standardized','g', num_ftrs + ord_ftrs)
plt.annotate('Correct!',xy=(10,10),xycoords = 'axes points',color='g')

plt.tight_layout()
plt.savefig('figures/coefs_comparison.png',dpi=300)
plt.show()

```



Introduction to Interpretability in ML

By the end of this workshop, you will be able to

- calculate global explanations using the coefficients of linear models
- **perform permutation importance using any ML model**
- use the SHAP package to calculate local explanations

Permutation feature importance

- model agnostic, you can use it with any supervised ML model
- steps:
 - train a model and calculate a test score :)
 - randomly shuffle a single feature in the test set
 - recalculate the test score with the shuffled data
 - model score worsens because the shuffling breaks the relationship between feature and target
 - the larger the difference, the more important the feature is

```
In [5]: from sklearn.svm import SVR
param_grid = {'regressor__C': [0.01, 0.1, 1, 10, 100, 1000, 10000],
              'regressor__gamma': [1e-4, 1e-3, 0.01, 0.1, 1, 10, 100]}

model, test_score, train_set, test_set = MLpipe_KFold_RMSE(X, y, preprocessor3,
print(test_score)

{'regressor__C': 1000, 'regressor__gamma': 0.01}
0.5645801016038898 0.5645801016038898
0.5645801016038898
```

```
In [6]: np.random.seed(42)

X_test = test_set[0]
```

```

y_test = test_set[1]
ftr_names = X_test.columns

nr_runs = 10
scores = np.zeros([len(ftr_names),nr_runs])

print('test score = ',test_score)
# loop through the features
for i in range(len(ftr_names)):
    print('shuffling '+str(ftr_names[i]))
    r2_scores = []
    for j in range(nr_runs):
        X_test_shuffled = X_test.copy()
        X_test_shuffled[ftr_names[i]] = np.random.permutation(X_test[ftr_names[i]])
        r2_scores.append(model.score(X_test_shuffled,y_test))
    print('    shuffled test score:',np.around(np.mean(r2_scores),3),'+/-',np.around(np.std(r2_scores),3))
    scores[i] = r2_scores

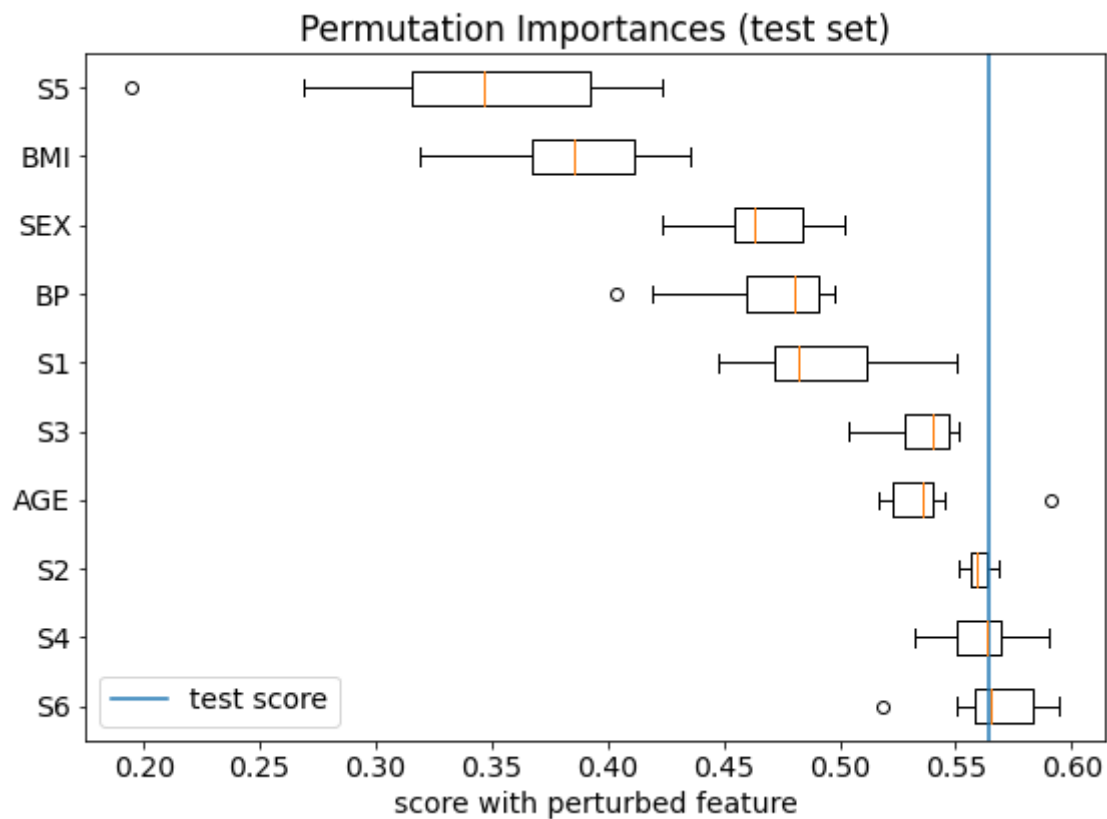
test score = 0.5645801016038898
shuffling AGE
    shuffled test score: 0.537 +/- 0.02
shuffling SEX
    shuffled test score: 0.467 +/- 0.021
shuffling BMI
    shuffled test score: 0.384 +/- 0.037
shuffling BP
    shuffled test score: 0.468 +/- 0.031
shuffling S1
    shuffled test score: 0.493 +/- 0.034
shuffling S2
    shuffled test score: 0.56 +/- 0.005
shuffling S3
    shuffled test score: 0.535 +/- 0.014
shuffling S4
    shuffled test score: 0.561 +/- 0.016
shuffling S5
    shuffled test score: 0.34 +/- 0.066
shuffling S6
    shuffled test score: 0.567 +/- 0.021

```

```

In [7]: sorted_indcs = np.argsort(np.mean(scores,axis=1))[:,::-1]
plt.rcParams.update({'font.size': 14})
plt.figure(figsize=(8,6))
plt.boxplot(scores[sorted_indcs].T,labels=ftr_names[sorted_indcs],vert=False)
plt.axvline(test_score,label='test score')
plt.title("Permutation Importances (test set)")
plt.xlabel('score with perturbed feature')
plt.legend()
plt.tight_layout()
plt.show()

```

Check out sklearn's permutation importance!

[https://scikit-](https://scikit-learn.org/stable/modules/generated/sklearn.inspection.permutation_importance.html)

[learn.org/stable/modules/generated/sklearn.inspection.permutation_importance.html](https://scikit-learn.org/stable/modules/generated/sklearn.inspection.permutation_importance.html)

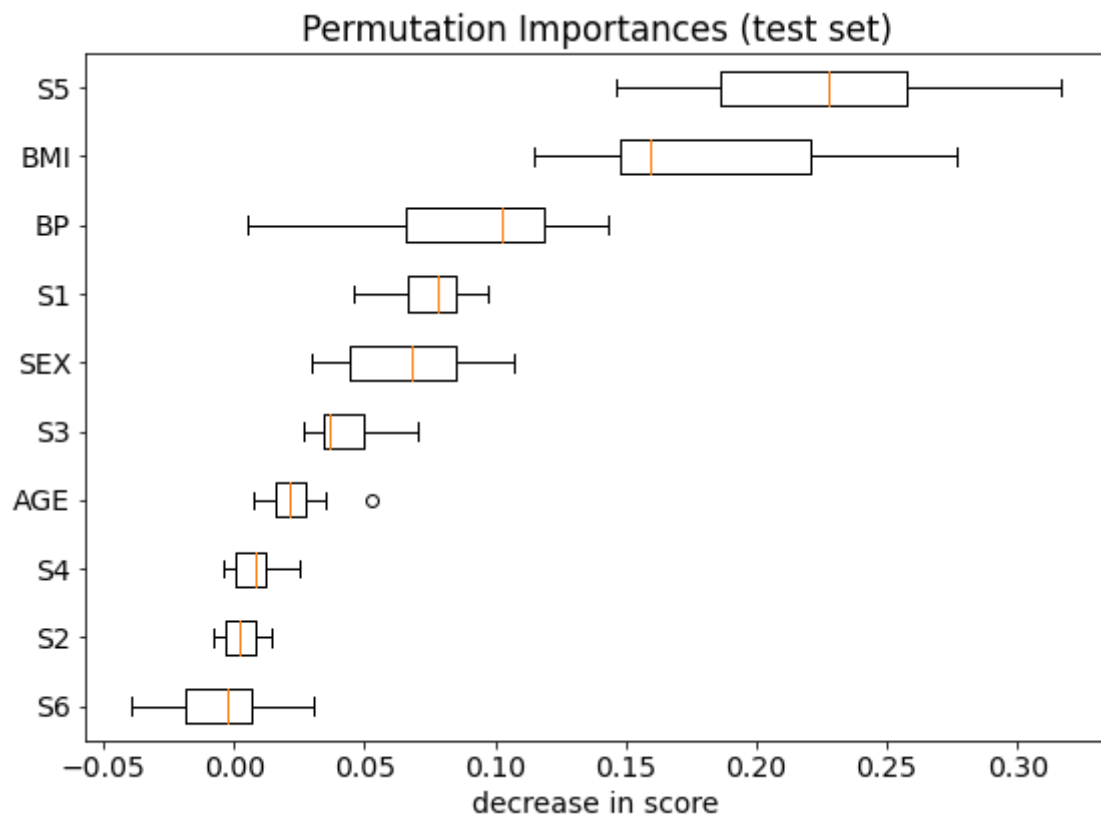
https://scikit-learn.org/stable/modules/permutation_importance.html#permutation-importance

```
In [8]: from sklearn.inspection import permutation_importance

result = permutation_importance(model, X_test, y_test, n_repeats=10, random_state=42)
ftr_names = X_test.columns

sorted_indcs = result.importances_mean.argsort()

plt.rcParams.update({'font.size': 14})
plt.figure(figsize=(8,6))
plt.boxplot(result.importances[sorted_indcs].T, labels=ftr_names[sorted_indcs], vert=False)
plt.title("Permutation Importances (test set)")
plt.xlabel('decrease in score')
plt.tight_layout()
plt.show()
```



Cons of permutation feature importance

- strongly correlated features
 - if one of the features is shuffled, the model can still use the other correlated feature
 - both features appear to be less important but they might actually be important
 - solution:
 - check the correlation matrix plot
 - remove all but one of the strongly correlated features
- no feature interactions
 - one feature might appear unimportant but combined with another feature could be important
 - solution:
 - permute two features to measure how important feature pairs are
 - this can be computationally expensive

Introduction to Interpretability in ML

By the end of this workshop, you will be able to

- calculate global explanations using the coefficients of linear models
- perform permutation importance using any ML model
- **use the SHAP package to calculate local explanations**

SHAP values

- one way to calculate local feature importances
- it is based on Shapely values from game theory
- read more [here](#), [here](#), and [here](#)

Cooperative game theory

- A set of m players in a coalition generate a surplus.
- Some players contribute more to the coalition than others (different bargaining powers).
- How important is each player to the coalition?
- How should the surplus be divided fairly amongst the players?

Cooperative game theory **applied to feature attribution**

- A set of m **features** in a **model** generate a **prediction**.
- Some **features** contribute more to the **model** than others (different **predictive** powers).
- How important is each **feature** to the **model**?
- How should the **prediction** be divided amongst the **features**?

How is it calculated?

$$\Phi_i = \sum_{S \subseteq M \setminus i} \frac{|S|!(M-|S|-1)!}{M!} [f_x(S \cup i) - f_x(S)]$$

- Φ_i - the contribution of feature i
- M - the number of features
- S - a set of features excluding i , a vector of 0s and 1s (0 if a feature is missing)
- $|S|$ - the number of features in S
- $f_x(S)$ - the prediction of the model with features S

How is it calculated?

$$\Phi_i = \sum_{S \subseteq M \setminus i} \frac{|S|!(M-|S|-1)!}{M!} [f_x(S \cup i) - f_x(S)]$$

- the difference feature i makes in the prediction:
 - $f_x(S \cup i)$ - the prediction with feature i
 - $f_x(S)$ - the prediction without feature i
 - loop through all possible ways a set of S features can be selected from the M features excluding i
 - **weight the contribution based on how many ways we can select $|S|$ features**
-

```
In [9]: import shap
shap.initjs() # required for visualizations later on

explainer = shap.KernelExplainer(model[1].predict,model[0].transform(train_set[
# transform the test set
X_test_transformed = model[0].transform(X_test)
print(np.shape(X_test_transformed))
# calculate shap values on the first 1000 points in the test
shap_values = explainer.shap_values(X_test_transformed)
print(np.shape(shap_values))
ftr_names = num_ftrs + ord_ftrs
```



Using 353 background data samples could cause slower run times. Consider using `shap.sample(data, K)` or `shap.kmeans(data, K)` to summarize the background as K samples.


```
(89, 10)
0%|          | 0/89 [00:00<?, ?it/s]
(89, 10)
```

Explain a point

```
In [10]: ftr_names = np.array(num_ftrs + ord_ftrs)

index = 1 # the index of the point to explain
shap.force_plot(explainer.expected_value, shap_values[index,:], X_test_transfor
```

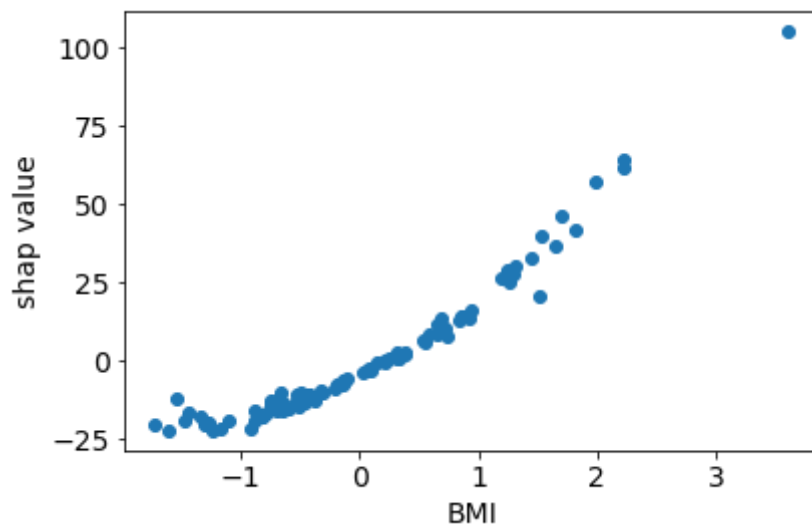
Out[10]:



Feature value vs. shap value


```
In [11]: ftr = 'BMI'
indx = np.argwhere(ftr_names==ftr)

plt.scatter(X_test_transformed[:,indx],shap_values[:,indx])
plt.ylabel('shap value')
plt.xlabel(ftr)
plt.tight_layout()
plt.show()
```



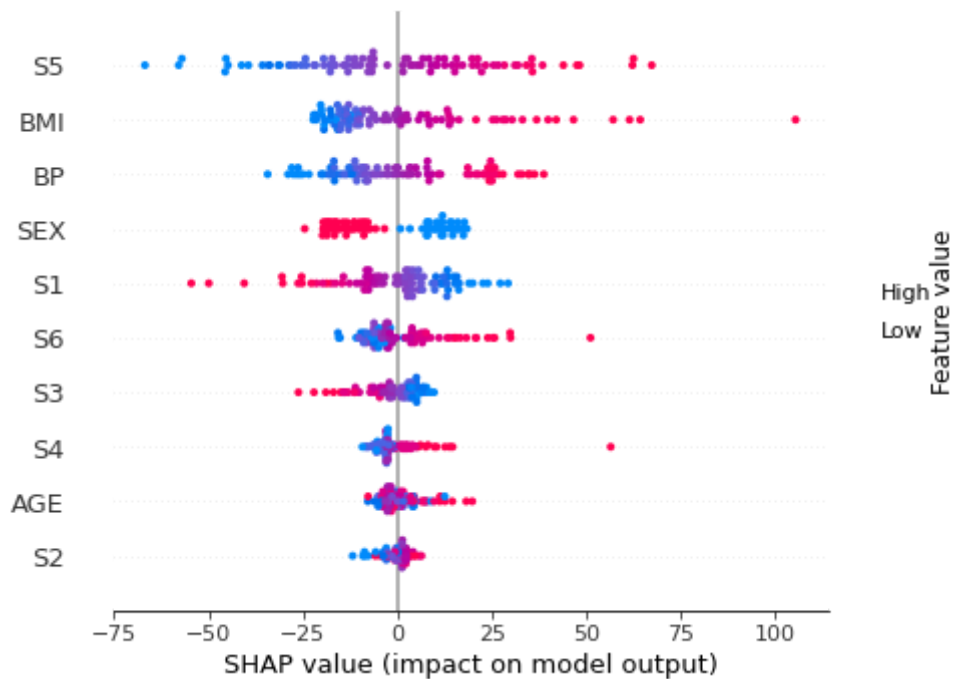
Dependence plot

In [12]: `shap.dependence_plot('BMI', shap_values, X_test_transformed, feature_names=ftr_`



It can also be used for global feature importance

In [13]: `shap.summary_plot(shap_values, X_test_transformed, feature_names = ftr_names)`



SHAP cons

- it can be numerically expensive
 - an efficient shap method was developed for trees, see [here](#)
- how to estimate $f_x(S)$?
 - this is not trivial because models cannot change the number of features they use
 - usually the values of the dropped features are replaced with the mean or 0
 - this is an approximation

Recap: Introduction to Interpretability in ML

By now, you can

- calculate global explanations using the coefficients of linear models
- perform permutation importance using any ML model
- use the SHAP package to calculate local explanations

Brown Data Science Master's Program

- if your company is looking to hire talented data scientists, reach out to me
- our students are looking for internships and full time positions!
- email: andras_zsom@brown.edu
- I'm also around tomorrow and until the early afternoon on Thursday